

PlantTech report

Carlos Leon

Elevator Pitch

This project aims to show the process of two things: the construction of a machine learning model to explain a client dataset, and the reduction the number of features the client will need to take care of.

Dataset insights

The file `assignment.csv` contains collected information of 24 features named as `par_0` , `par_1` , ..., and the target `y` . Missing values and normalization procedures can be applied through a pipeline.




Fig.1 Scatter plot of the first 3 features and the target. Strong colors shows a higher density

distribution. Diagonal plots display histograms

Fig.1 exhibits some correlation between `par_0` and `y` , but poor correlation between `par_` features (a scatter plot between all the targets is attached, see `scatter.png`). Therefore, features can be approximated as independent parameters. A more detailed answer on the correlation between features and target is shown in Table 1.

	correlation
y	1.00
par_4	0.32
par_12	0.21
par_8	0.12
par_16	0.10
par_0	0.07
par_20	0.06
par_11	0.03

Table 1. Absolute values of the correlation between target and parameters

From Table 1, `par_4`, `12`, `8`, `16` and `0` are the five features that contribute the most to the target. The entire correlation matrix is shown in the file `runMe.ipynb` and it shows the feature independence mentioned before (very low values of correlation between parameters).

Machine learning models

Fine-tune model

After setting aside a validation set (splitting of 20%), linear regressions and random forest regressions has been applied to the entire training set with 24 features, and also to a subset containing only the best 5 features.

RMSE values

	24 features	5 best features
--	-------------	-----------------

	24 features	5 best features
linear terms	0.046	0.047
polynomial terms	--	0.015
random forest	0.007	0.006

Table 2. RMSE values for the chosen ML models. Linear regressor models were tested with manipulated input containing only linear terms or 5th-degree polynomial terms.

To obtain the RMSE values shown in Table 2, models were tested for a small subset of the training set (a more complex procedure would involve cross-validation). This process allows us to choose the best model for which the validation test will be employed: random forest (expensive: 2 min) and/or linear regressor for 5th-degree polynomials (cheap: seconds). Default scikit-learn parameters were used for the random forest regressors. Here, a grid search cross-validation was also performed, but without significant improvement on the RMSE value.

Validation

Finally, the chosen models were applied to predict the targets on the validation test. Table 3 shows no significant different in the RMSE value between the two models. However, the linear regressor using polynomial terms was computationally cheaper.

Fig. 2 shows the relationship between predicted and target values. An ideal model would locate this relationship over the identity function (black line).

RMSE values

	5 best features
polynomial terms	0.0286
random forest	0.0295

Table 3. Linear regressor with polinomial terms vs random forest RMSE values for the validation test




Fig 2. Prediction vs target relationship for the validation test

Conclusions

A linear regressor model and a 5th-degree polynomial combinations for the 5 best features can deliver the same RMSE values than other complex models studied here, and is able to employ less computational resources. The client may focus on the following features: `par_4` , `par_12` , `par_8` , `par_16` , `par_0` .